Estimation of marginal effects for models with alternative variable transformations

Fernando Rios-Avila
Levy Economics Institute
Annandale-on-Hudson, NY
friosavi@levy.org

Abstract. `margins` is a powerful postestimation command that allows the estimation of marginal effects for official and community-contributed commands, with well-defined predicted outcomes (see `predict`). While the use of factor-variable notation allows one to easily estimate marginal effects when interactions and polynomials are used, estimation of marginal effects when other types of transformations such as splines, logs, or fractional polynomials are used remains a challenge. In this article, I describe how `margins`'s capabilities can be extended to analyze other variable transformations using the command `f_able`.

Keywords: st0628, f_able, margins, marginal effects, predict, variable transformations, nonlinear

1 Introduction

`margins` is a powerful postestimation command that was introduced in Stata 11, allowing the estimation of marginal effects for all official estimation commands and any community-contributed command with a properly defined predicted outcome program (see `predict`). As described in Williams (2012), `margins`, and its companion `marginsplot` (in Stata 12), provides a great tool for analyzing and providing meaningful information that is easier to interpret and describe. Furthermore, in combination with factor-variable notation, also introduced in Stata 11, `margins` also allows users to easily estimate marginal effects when interactions and polynomials of continuous variables, and interactions with discreet variables, are used without any additional work from the user.

Despite these features, the ability of `margins` to estimate marginal effects is limited to the use of simple polynomials and interactions. Unless further information is provided to a command, Stata cannot identify the possible interdependencies across explanatory variables.

To the best of my knowledge, there are only two commands where this limitation is not binding. As described in Poi (2011), `nl` can be used to obtain marginal effects when using functions other than interactions of polynomials, even when the model is nonlinear in parameters, but assuming an additive error. `npregress series`, one of the newest additions to Stata 16 nonparametric analysis, also estimates marginal effects of arbitrary transformations of the independent variables (splines, B-splines, and polynomials), based on numerical derivatives. Unfortunately, both models are based on

least-squares types of estimators. Neither command can be used in the case of nonlinear models like logit or probit. In the case of `npregress series`, users cannot freely choose what transformations to use. The only other program that proposes a strategy that may allow obtaining marginal effects, more specifically marginal means, using transformed covariates, was proposed by Royston (2013). His command, `marginscontplot`, uses a predefined mapping between original and transformed covariates to appropriately handle the constructed variables in `margins`.

This article describes a simple strategy, implemented by `f_able`, that allows the estimation of marginal effects of a larger set of nonlinear transformations using `margins` and the option `nochainrule [numerical noestimcheck]`. Section 2 provides a review of how marginal effects should be estimated, with emphasis on some of the computational challenges that are bypassed by the use of factor notation. Section 3 provides a simple approach for the estimation of marginal effects for the case of Spline regressions and compares it with a simplified case using `npregress series`. Section 4 introduces three commands that prepare the data for the estimation of marginal effects for any arbitrary variable transformation and shows how `f_able` can be used to estimate marginal effects using any variable transformation. Section 5 concludes.

2   Marginal effects: Theoretical approach

Marginal effects are useful statistics to measure the impact of a change that the independent variables will have on the dependent variable, assuming other covariates remain constant. At the beginning of most introductory econometric courses, very little emphasis is placed on understanding this concept because, for linear regressions, marginal effects are typically equal to the coefficient associated with the variable analyzed. Consider the following linear regression model:

$$y_i = b_0 + b_1 x_{1i} + b_2 x_{2i} + e_i \tag{1}$$

Under the standard assumptions of exogeneity, homoskedasticity, and correct model specifications (see Wooldridge [2013]), the coefficients of (1) can be estimated using ordinary least squares (OLS). Under the ceteris paribus assumption, the marginal effect of x_1 and x_2 on y could be estimated based on the thought experiment of what would the outcome y_i be if x_{1i} (x_{2i}) increases in one unit, keeping everything else constant. Mathematically, this is equivalent to obtaining the partial derivative of y_i with respect to x_{1i}, which for (1) is given by b_1 (b_2):

$$\Delta y_i = \Delta x_i b_1 \rightarrow \frac{\Delta y_i}{\Delta x_{1i}} = \frac{\partial y_i}{\partial x_{1i}} = b_1; \quad \frac{\partial y_i}{\partial x_{2i}} = b_2$$

However, if we lift the homoskedasticity assumption, it is more convenient to estimate marginal effects using the conditional expectation assumption. Under this assumption, we estimate the average effect of a one-unit change in x_1 or x_2 across all individuals with the same observed characteristics (conditioning on x):

$$E(y_i|x) = b_0 + b_1 x_1 + b_2 x_2$$

$$\frac{\partial E(y_i|x)}{\partial x_1} = b_1; \quad \frac{\partial E(y|x)}{\partial x_2} = b_2$$

This implies that for a simple linear regression [like (1)], where each variable appears only once, and without any transformation, marginal effects are directly identified by the estimated coefficients.

Later in a semester, students are introduced to the idea that nonlinear transformations and interactions of the independent variables can be included in the linear regression model. Because the model is still linear in parameters, it can be estimated using OLS. However, additional care needs to be taken when estimating the marginal effects because one needs to account for the interdependence of variable transformations. Consider the following model, with its conditional expectation form:

$$y_i = b_0 + b_1 x_{1i} + b_2 x_{1i}^2 + b_3 x_{2i} + b_4 x_{1i} x_{2i} + e_i$$
$$E(y_i|x) = b_0 + b_1 x_1 + b_2 x_1^2 + b_3 x_2 + b_4 x_1 x_2 \tag{2}$$

In this model, the marginal effect of x_1 on y is no longer a constant, and it now depends on the values x_1 and also of x_2. Using calculus, the marginal effects for x_1 and x_2 in this model are

$$\frac{\partial E(y_i|x)}{\partial x_1} = b_1 + 2b_2 x_1 + b_4 x_2; \quad \frac{\partial E(y_i|x)}{\partial x_2} = b_3 + b_4 x_1$$

Because marginal effects are now functions of x_1 and x_2, one usually needs to decide what values of x_1 and x_2 to use when reporting marginal effect. Alternatively, one can also estimate marginal effects for all observed (or feasible) combinations of x_1 and x_2 and summarize (or plot) those estimates as needed.[1] The standard practice is presenting average marginal effects (AME) or the marginal effects at the mean (MEM). For (2), they both will be the same:

AME   MEM

$$E\left(\frac{\partial E(y_i|x)}{\partial x_1}\right) \quad \left.\frac{\partial E(y_i|x)}{\partial x_1}\right|_{x=\bar{x}} = b_1 + 2b_2\bar{x}_1 + b_4\bar{x}_2$$

$$E\left(\frac{\partial E(y_i|x)}{\partial x_2}\right) \quad \left.\frac{\partial E(y_i|x)}{\partial x_2}\right|_{x=\bar{x}} = b_3 + b_4\bar{x}_1$$

With this information, and assuming x_1 and x_2 are nonstochastic, standard errors associated with these marginal effects can be estimated,[2] and the technical part of the analysis would be done.

1. Note that the command `margins` can generate marginal effects for all available data using the undocumented option `gen()`. I want to thank an anonymous referee for pointing this out.
2. A detailed explanation of how standard errors are estimated using `margins` can be found at https://www.stata.com/support/faqs/statistics/compute-standard-errors-with-margins/.

The challenge with most software (including Stata) is that, unless additional information is provided, the software may not recognize that some variables are interrelated by construction and that the assumption of "everything else remaining constant" is incorrect. In the case of Stata, until factor variables were introduced, it could not automatically adjust for these interrelations, providing incorrect estimations of marginal effects, unless further steps were considered. Stata still fails to grasp these interrelations when we step beyond simple interactions or polynomials.

3   Marginal effects: Empirical approach

For this section, `dui.dta`, which is freely available online will be used. Consider the following model:

$$\texttt{citations}_i = b_0 + b_1 \texttt{fines}_i + b_2 \texttt{fines}_i^2 + e_i \tag{3}$$

Before Stata 11 and factor notation, to fit a model like this, one would need to create all variables before including them in the model. For example, one would need to create a variable named `fines2` to be equal to $\texttt{fines}_i^2$,

```
. webuse dui
(Fictional data on monthly drunk driving citations)
. generate fines2 = fines^2
```

and fit the model using the command `regress`:

```
. regress citations fines fines2
```

Source	SS	df	MS			
Model	20750.38	2	10375.19	Number of obs	=	500
Residual	27200.458	497	54.7292917	F(2, 497)	=	189.57
				Prob > F	=	0.0000
				R-squared	=	0.4327
Total	47950.838	499	96.0938637	Adj R-squared	=	0.4305
				Root MSE	=	7.3979

citations	Coef.	Std. Err.	t	P>|t|	[95% Conf. Interval]	
fines	-47.10883	7.691611	-6.12	0.000	-62.22091	-31.99674
fines2	1.98084	.3885844	5.10	0.000	1.21737	2.744311
_cons	293.0067	37.94286	7.72	0.000	218.4585	367.5549

Using `margins` to calculate the marginal effects of `fines` on `citations` would produce the wrong answer because `fines2` would not be handled correctly: there is no way for Stata to know that `fines2` = `fines`^{2}.

```
. margins, dydx(fines)
Average marginal effects                         Number of obs   =       500
Model VCE     : OLS

Expression    : Linear prediction, predict()
dy/dx w.r.t. : fines
```

	dy/dx	Delta-method Std. Err.	t	P>|t|	[95% Conf. Interval]	
fines	-47.10883	7.691611	-6.12	0.000	-62.22091	-31.99674

However, AME can be derived by hand by taking the derivative of (3) with respect to `fines`, using the average value of `fines` as the point of reference, and using `lincom` to calculate the marginal effects and standard errors:

```
. summarize fines, meanonly
. lincom _b[fines]+2*_b[fines2]*`r(mean)'
 ( 1)  fines + 19.7904*fines2 = 0
```

citations	Coef.	Std. Err.	t	P>|t|	[95% Conf. Interval]	
(1)	-7.907201	.4236816	-18.66	0.000	-8.739629	-7.074773

Since the introduction of Stata 11, estimating marginal effects for a model like this is much simpler. Using factor notation, we add the squared parameter and let `margins` handle the interaction on its own:

```
. quietly regress citations fines c.fines#c.fines
. margins, dydx(fines)
Average marginal effects                         Number of obs   =       500
Model VCE     : OLS

Expression    : Linear prediction, predict()
dy/dx w.r.t. : fines
```

	dy/dx	Delta-method Std. Err.	t	P>|t|	[95% Conf. Interval]	
fines	-7.907201	.4236816	-18.66	0.000	-8.739629	-7.074773

I suspect that the methodology behind the code for `margins` and factor notation "taught" Stata how to take derivatives of polynomials. In other words, Stata recognizes that when there is an expression like `c.var1#c.var1`, it internally "knows" the analytical derivative is $2 \times$ `c.var1`. Thus, `margins` simply uses this information to handle the squared parameter (`c.fines#c.fines`) before providing you with the result.

While this has been a big improvement for a better understanding of marginal effects, it does have its limitations. For example, `margins` would not be able to estimate marginal effects for the following models:

$$\text{Model 1}: \texttt{citations}_i = b_0 + b_1(1/\texttt{fines}_i) + e_i \tag{4}$$
$$\text{Model 2}: \texttt{citations}_i = b_0 + b_1\texttt{fines}_i^{0.5} + e_i$$
$$\text{Model 3}: \texttt{citations}_i = b_0 + b_1\texttt{fines} + b_2\max(\texttt{fines} - 9.9, 0) + e_i$$

However, mathematically, the AME, and MEM, can be derived directly using the sample means:

	AME	MEM
Model 1 :	$b_1 n^{-1} \sum \left(-\dfrac{1}{\texttt{fines}_i^2}\right)$	$-b_1 \dfrac{1}{(n^{-1}\sum\texttt{fines}_i)^2}$
Model 2 :	$0.5 \times b_1 n^{-1} \sum \left(\texttt{fines}_i^{-0.5}\right)$	$0.5 \times b_1 \left(n^{-1}\sum\texttt{fines}_i\right)^{-0.5}$
Model 3 :	$b_1 + b_2 n^{-1} \sum \{1(\texttt{fines}_i > 9.9)\}$	$b_1 + b_2 \times 1\left\{\left(n^{-1}\sum\texttt{fines}_i\right) > 9.9\right\}$

Which can be used to estimate the AME by hand. Now, concentrate on the estimation of the AME:

```
. webuse dui, clear
(Fictional data on monthly drunk driving citations)
. generate i_fines = 1/fines
. generate ni_fines2 = -1/fines^2
. generate fines05 = fines^.5
. generate i_fines05 = 0.5*fines^-.5
. generate fines_99 = max((fines-9.9),0)
. generate dfines_99 = fines>9.9
. * model 1
. quietly regress citations i_fines
. summarize ni_fines2, meanonly
. lincom _b[i_fines]*`r(mean)'
 ( 1)  - .0104108*i_fines = 0
```

citations	Coef.	Std. Err.	t	P>|t|	[95% Conf. Interval]
(1)	-8.091099	.4224389	-19.15	0.000	-8.921081 -7.261117

```
. * model 2
. quietly regress citations fines05
. summarize i_fines05, meanonly
```

```
. lincom _b[fines05]*`r(mean)´
 ( 1)   .1593264*fines05 = 0
```

citations	Coef.	Std. Err.	t	P>|t|	[95% Conf. Interval]
(1)	-8.010351	.4314167	-18.57	0.000	-8.857972 -7.162729

```
. * model 3
. quietly regress citations fines fines_99
. summarize dfines_99, meanonly
. lincom _b[fines]+_b[fines_99]*`r(mean)´
 ( 1)   fines + .5*fines_99 = 0
```

citations	Coef.	Std. Err.	t	P>|t|	[95% Conf. Interval]
(1)	-7.926694	.4271729	-18.56	0.000	-8.765981 -7.087407

For these models, one also has the option of using `nl` to estimate the marginal effects, which requires less work:[3]

```
. quietly nl (citations = {b0}+{b1}/fines), variable(fines)
. margins, dydx(fines)
  (output omitted)
. quietly nl (citations = {b0}+{b1}*fines^.5), variable(fines)
. margins, dydx(fines)
  (output omitted)
. quietly nl (citations = {b0}+{b1}*fines+{b2}*max((fines-9.9),0)),
> variable(fines)
. margins, dydx(fines)
  (output omitted)
```

This flexibility of `nl` to estimate marginal effects with nonstandard transformations motivates the following question: why can `nl` "correctly" estimate marginal effects, whereas `regress` cannot? The answer is rather simple: we are not using a constructed variable in the model; instead, we are using the original variable and allowing `nl` to construct the new variable.

This means that, while we see this model is being fit,

```
citations = {b0}+{b1}*fines^.5
```

what may be happening in the background is different. Stata identifies which elements of this code are parameters to be estimated (those within brackets) and which elements need to be created (`fines`$^{.5}$) before fitting the model. In other words, Stata is estimating the following:

```
citations = {b0}+{b1}*__000000
```

---

3. For this to work, we need to indicate to `nl` that `fines` is a variable in the model. Outputs are identical to the ones produced by hand.

Here `__000000` is a temporary variable that was constructed as $\texttt{fines}^{.5}$ and that you never see. The difference is that `nl` knows that $\texttt{\_\_000000} = \texttt{fines}^{.5}$.

How is it that `nl` knows that the derivative of $\texttt{fines}^{0.5}$ is $0.5\texttt{fines}^{-0.5}$? The answer is `nl` does not know. The only type of analytical derivatives Stata may know is when interactions are used, specifically factor notations. However, because Stata "remembers" how a variable was constructed, it can use numerical derivatives to make a reasonable approximation for the analytical derivative.

For example, the numerical derivative for the transformation $\texttt{fines}^{0.5}$ can be approximated as follows:

$$\frac{\partial(\texttt{fines}^{0.5})}{\partial\texttt{fines}} = \frac{(\texttt{fines}+h)^{.5} - (\texttt{fines}-h)^{.5}}{2h} \quad \text{for a sufficiently small } h$$

This expression is surprisingly accurate.[4] For this example, when $h = 1$, the largest absolute difference between the numerical and analytical derivative is 0.000423, whereas when $h = 1/216$, the largest difference is 6.58e–09. This implies that `margins` does not need to know how to obtain analytical derivatives, because it can use numerical derivatives instead and use this information to estimate the appropriate marginal effects.

4   f_able: Marginal effects for arbitrary variable transformations

The previous sections described how marginal effects should be estimated and how to use that information to obtain AME or MEM by comparing the step-by-step procedure with what `margins` does. They also detailed how `nl` can estimate marginal effects for variable transformations other than interactions using numerical derivatives to approximate the analytical derivatives. This implies that Stata already has the capability to estimate marginal effects when transformations other than interactions and polynomials are used in a model. However, to expand those capabilities to any linear and nonlinear model, we need to "tell" Stata that some variables are based on constructions from others so that Stata can use this information to correctly estimate marginal effects.

In this section, I first sketch the basic structure of the programs required to solve this problem and then describe the commands within `f_able` that address this problem.

4.1   f_able: Basics

One solution to pass information about variable dependencies to Stata is to use the variable label to "store" the transformation used to create the variable of interest. To automate this procedure, I propose two small programs that "wrap" around Stata's built-in commands `generate` and `replace`:

---

4. In fact, as indicated in its help file, the function `deriv()` in Mata uses this approximation to compute numerical derivatives.

```
. program fgen
  1.     syntax newvarname =/exp [if] [in]
  2.     generate `typelist´ `varlist´=`exp´
  3.     label var `varlist´ "`exp´"
  4. end
. program frep
  1.     syntax varname =/exp [if] [in]
  2.     replace   `varlist´=`exp´
  3.     label var `varlist´ "`exp´"
  4. end
```

These two commands have one purpose: when a new variable is created, it will label it with the expression used to create it, and if the values are replaced, it will change the label to the new expression used.

Consider model 3 from (4). The variable of interest can be created using the command fgen, but I created it with an intentional mistake:

```
. fgen fines2 = max(fines-9,0)
. describe fines2

              storage   display    value
variable name  type     format     label       variable label
```

```
fines2         float    %9.0g                   max(fines-9,0)
```

However, I can rectify it using the command frep and replace the values fines2 with the correct content:

```
. frep fines2 = max(fines-9.9,0)
(420 real changes made)
. describe fines2

              storage   display    value
variable name  type     format     label       variable label
```

```
fines2         float    %9.0g                   max(fines-9.9,0)
```

In addition to setting up the data for the next step of the estimation of marginal effects, these commands may be useful for tracking how variables are created or modified.

The second step is to explicitly tell Stata that a particular variable is constructed based on other variables in the model and that when margins is called, the constructed variable needs to be updated every time the original variable changes. This can be done with the two following programs:

```
. program f_able, eclass
  1.      syntax, [* NLvar(varlist)]
  2.      _ms_dydx_parse `nlvar´
  3.      if "`e(predict_old)´"=="" {
  4.          ereturn local predict_old  `e(predict)´
  5.          ereturn local predict  f_able_p
  6.      }
  7.      foreach i of varlist `nlvar´ {
  8.          local fnc:variable label `i´
  9.          ereturn hidden local _`i´ `fnc´
 10.      }
 11.      ereturn local nldepvar `nlvar´
 12. end
. program f_able_p
  1.      syntax newvarname [if] [in], [*]
  2.      local idepvar `e(nldepvar)´
  3.      foreach i of local idepvar {
  4.          tempvar _`i´
  5.          quietly clonevar `_`i´´=`i´
  6.          quietly recast double `i´
  7.          quietly replace `i´=`e(_`i´)´
  8.      }
  9.      `e(predict_old)´ `0´
 10.      foreach i of local idepvar {
 11.          if "`i´"!="_cons" {
 12.              quietly replace `i´=`_`i´´
 13.          }
 14.      }
 15. end
```

The first program, f_able, does three things. First, it adds information to any previously fit model indicating which variables are constructed variables using e(nldepvar). Second, it adds hidden macros with information regarding the data transformation used.[5] Last, it redirects predict from the original e(predict) command to the one defined below as f_able_p but keeps the original information in e(predict_old).

The second program, f_able_p, has only one purpose—updating the constructed variables identified in e(nldepvar)—before obtaining the predicted values appropriate for the estimated command e(predict_old), using the information previously stored in the hidden macros.

With these two pieces of code, the last step is to call margins for the estimation of the marginal effects, using the option nochainrule.[6] This is an option often overlooked when using official Stata commands. However, as its help file indicates, "nochainrule is safer because it makes no assumptions about how the parameters and covariates join to form the response."

---

5. While this step is unnecessary in most cases I have considered, in some examples, margins "drops" the variable label information. In those cases, storing the information as part of the estimation commands seems to be a better alternative.

6. Some commands, such as logit, probit, and poisson also require one to include the option numerical to use numerical derivatives.

This implies that when a model like the following is fit,

```
. regress citations fines fines2
  (output omitted)
. f_able, nl(fines2)
```

marginal effects will be calculated using the coefficient of `fines` and the coefficient of `fines2` times the numerical derivative of `fines2` with respect to `fines`, if `fines2` was declared as a constructed variable. Let see how this works.

```
. quietly regress citations fines fines2
. f_able, nlvar(fines2)
. margins, dydx(fines) nochainrule
```

Average marginal effects   Number of obs   =   500
Model VCE   : OLS
Expression   : Fitted values, predict()
dy/dx w.r.t. : fines

	dy/dx	Delta-method Std. Err.	z	P>|z|	[95% Conf. Interval]
fines	−7.926694	.4271729	−18.56	0.000	−8.763937 −7.089451

The first line fits the model with `citations` as a dependent variable and `fines` and `fines2` as independent variables. `fines2` is the one we defined previously as $\max(\text{fines} − 9.9, 0)$. The second line calls `f_able` to identify that `fines2` is a constructed variable, and the last step estimates the marginal effects using `nochainrule`. The results are the same as those done by hand or those using the `nl` command. They can also be replicated using `npregress series`.

```
. npregress series citations fines, spline(1) knots(1)
Computing approximating function

Computing average derivatives
```

Linear-spline estimation   Number of obs   =   500
                                                  Number of knots =   1

citations	Effect	Robust Std. Err.	z	P>|z|	[95% Conf. Interval]
fines	−7.926694	.4772213	−16.61	0.000	−8.86203 −6.991357

Note: Effect estimates are averages of derivatives.

Notice that while the point estimates are the same, the standard errors produced by `npregress series` are somewhat larger than those produced with `nl` or with the proposed strategy, because `npregress` reports robust standard errors by default.

We can also compare the output of `f_able` with the output we could obtain using factor notation:

```
. quietly regress citations c.fines##c.fines##c.fines

. margins, dydx(fines)
Average marginal effects                          Number of obs     =        500
Model VCE     : OLS

Expression    : Linear prediction, predict()
dy/dx w.r.t. : fines
```

		Delta-method				
	dy/dx	Std. Err.	t	P>|t|	[95% Conf. Interval]	
fines	-7.928817	.4226225	-18.76	0.000	-8.759168	-7.098465

```
. frep fines2 = fines^2
(500 real changes made)

. fgen fines3 = fines^3

. quietly regress citations fines fines2 fines3

. f_able, nlvar(fines2 fines3)

. margins, dydx(fines) nochainrule
Average marginal effects                          Number of obs     =        500
Model VCE     : OLS

Expression    : Fitted values, predict()
dy/dx w.r.t. : fines
```

		Delta-method				
	dy/dx	Std. Err.	z	P>|z|	[95% Conf. Interval]	
fines	-7.928819	.4226228	-18.76	0.000	-8.757144	-7.100493

This method replicates the results when using factor notation up to five decimal places, but some degree of precision is lost because of forced use of numerical derivatives.

Last, here is a more challenging estimation that combines the use of factor notation with f_able and compares the output with that of the npregress series.

```
. npregress series citations fines i.csize, spline(2) knots(1)
Computing approximating function

Computing average derivatives
Quadratic-spline estimation                    Number of obs     =          500
                                               Number of knots   =            1
```

citations	Effect	Robust Std. Err.	z	P>|z|	[95% Conf. Interval]	
fines	-7.590817	.3533786	-21.48	0.000	-8.283427	-6.898208
csize (medium vs small)	5.48074	.5827045	9.41	0.000	4.33866	6.622819
(large vs small)	10.69879	.6043375	17.70	0.000	9.514311	11.88327

```
Note: Effect estimates are averages of derivatives for continuous covariates
      and averages of contrasts for factor covariates.
```

The same can be replicated using regress and f_able:

```
. webuse dui, clear
(Fictional data on monthly drunk driving citations)
. fgen double fines2 = fines^2
. fgen double fines3 = max(fines-9.9,0)^2
. quietly regress citations c.(fines fines2 fines3)##i.csize
. f_able, nlvar(fines2 fines3)
. margins, dydx(fines csize) nochainrule
Average marginal effects                       Number of obs     =          500
Model VCE    : OLS

Expression   : Fitted values, predict()
dy/dx w.r.t. : fines 2.csize 3.csize
```

	dy/dx	Delta-method Std. Err.	z	P>|z|	[95% Conf. Interval]	
fines	-7.590817	.3350832	-22.65	0.000	-8.247568	-6.934066
csize						
medium	5.480738	.6522748	8.40	0.000	4.202303	6.759174
large	10.69879	.6171921	17.33	0.000	9.489118	11.90847

```
Note: dy/dx for factor levels is the discrete change from the base level.
```

Once again, this replicates the output with very small differences in the point estimates but with some differences in terms of the standard errors, because npregress reports robust standard errors.

We can also use this strategy for models other than OLS. Consider `mksp2.dta`. As with the example provided in the help file for the command `mkspline`, four variables are created to allow for linear splines with four knots. A logit model of the outcome against `dosage` is fit, and the constructed variables are estimated. `margins` needs to include the options `nochainrule` and `numerical`. The predicted probabilities and the marginal effects across various values of `dosage` are produced:

```
. webuse mksp2, clear
. fgen dos1 = max(dosage-17.5,0)
. fgen dos2 = max(dosage-36.5,0)
. fgen dos3 = max(dosage-55.5,0)
. fgen dos4 = max(dosage-81.5,0)
. quietly logit outcome dosage dos1 dos2 dos3 dos4
. f_able, nl(dos1 dos2 dos3 dos4)
. quietly margins, nochainrule numerical at(dosage=(0(2)100))
. marginsplot, name(m1)
  Variables that uniquely identify margins: dosage
. quietly margins, dydx(dosage) nochainrule numerical at(dosage=(0(2)100))
. marginsplot, name(m2)
  Variables that uniquely identify margins: dosage
. graph combine m1 m2, xsize(8) scale(1)
```


Figure 1. Predicted probabilities and marginal effects

4.2   f_able package

In the previous subsection, I sketched the basic code that can be used to store information in regression outputs so that it can be used by `margins` to estimate marginal effects after most linear and nonlinear models. In this subsection, I describe all the programs that are part of `f_able` and can be used to address some limitations of the strategy.

```
fgen newvar = exp
frep oldvar = exp
```

`fgen` and `frep` can be used to create a new variable or replace an existing variable using any expression that is valid with `generate`. `fgen` creates a variable in double

format to reduce problems of precision with the constructed variables. If the expression is longer than 75 characters, it assigns the label `See note`, and a variable `note` will store the used expression.

    f_able, nlvar(*varlist*)

`f_able` is used to declare which variables in the model specification should be considered as "constructed" variables. All variables in `nlvar()`, which should be present in the previous model specification, are stored in a macro named `e(nldepvar)`. After one uses `f_able`, marginal effects with respect to the constructed variables are no longer possible unless `f_able_reset` is used. A list of hidden macros named `e(_varname)` is also stored and contains the information of how each variable (*varname*) was created.

`f_able_reset` restores all the postestimation results to the state previous to using `f_able`.

`f_symev` and `f_symrv` are commands that force the variance–covariance matrices stored in `e(V)` and `r(V)`, respectively, to be symmetric when a nonsymmetric warning appears. The warning appears because of the precision loss and forced numerical derivatives, which creates variance–covariance matrices that are nonsymmetric. This often occurs after using `margins` for multiple points of reference.

5   Conclusion

This article describes how marginal effects can be estimated using analytical derivatives, as well as numerical derivatives. It introduced two small programs, within the `f_able` package, that enable `margins` to estimate marginal effects when using transformations beyond variable interactions and polynomials.

These commands can be used to estimate marginal effects for many official Stata commands as well as other community-contributed commands that can produce sensible predicted outcomes. The strategy does have three limitations: first, the estimated marginal effects depend on the precision of the forced numerical derivatives; second, it requires the original variable to be present in the model specification so that marginal effects can be used; and third, because `f_able` is forcing `margins` to do something it is not meant to do, one may experience difficulties estimating marginal effects and standard errors when the request is particularly complex.

While the limitation regarding the precision of the estimates is unavoidable, the other limitations can be circumvented to some extent. In the second limitation, the original variable can be added to the list of explanatory variables using `o.`. This omits the original variable from the estimation but keeps it in the list of explanatory variables, allowing the estimation of `margins`.

The last limitation can be addressed with careful troubleshooting. While the option `nochainrule` forces `margins` to use numerical derivatives for the estimation of marginal effects after most commands, some commands such as `probit`, `logit`, and `poisson` require using the options `nochainrule` and `numerical`. If point estimates appear as

missing, the option `noestimcheck` can be used to bypass some of the safety checks in `margins` that produce this error. Finally, if standard errors are missing when estimating marginal effects for multiple points of interest, it may be because the estimated variance–covariance matrix created after `margins` is nearly symmetric. For those cases, `f_symev` and `f_symrv` can be used to force symmetry on the variance–covariance matrix produced by `margins`, so one can report standard errors or plot confidence intervals after `marginsplot`.

6   Programs and supplemental materials

To install a snapshot of the corresponding software files as they existed at the time of publication of this article, type

```
. net sj 21-1
. net install st0628      (to install program files, if available)
. net get st0628          (to install ancillary files, if available)
```

7   References

Poi, B. P. 2011. Stata tip 58: nl is not just for nonlinear models. *Stata Journal* 8: 139–141. https://doi.org/10.1177/1536867X0800800112.

Royston, P. 2013. marginscontplot: Plotting the marginal effects of continuous predictors. *Stata Journal* 13: 510–527. https://doi.org/10.1177/1536867X1301300305.

Williams, R. 2012. Using the margins command to estimate and interpret adjusted predictions and marginal effects. *Stata Journal* 12: 308–331. https://doi.org/10.1177/1536867X1201200209.

Wooldridge, J. M. 2013. *Introductory Econometrics: A Modern Approach*. 5th ed. Mason, OH: South-Western.

About the author

Fernando Rios-Avila is a research scholar at Levy Economics Institute of Bard College under the Distribution of Income and Wealth program. His research interests include applied econometrics, labor economics, and poverty and inequality.